



Individual Prefill

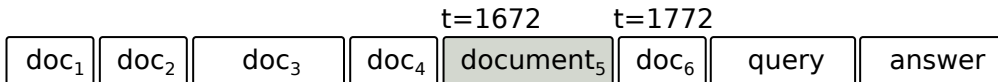
An empty prefix
or a BOS token



Cache reuse (inference)



RAG task



Recomputation methods

EPIC (): recompute the first K tokens




CacheBlend (): recompute K tokens that cause largest cache diff in the 1st layer



Fine tuning + Link token methods

BlockAttention (): Fine-tune the model to produce a good independent cache



KVLink (): in addition, train additional tokens and recover cross-attentions



KVPacket (): Link tokens only; head+tail; prefilled offline (not during inference)

